


DEEP LEARNING TUTORIAL

MAZIAR RAISSI

1. Outline. This is a short tutorial on the following topics in Deep Learning: Neural Networks, Recurrent Neural Networks, Long Short Term Memory Networks, Variational Auto-encoders, and Conditional Variational Auto-encoders. The full code for this tutorial can be found in <https://github.com/maziarraissi/DeepLearningTutorial>.

2. Neural Networks [2]. Consider the following deep neural network with two hidden layers.



Here, $\mathbf{x}_p \in \mathbb{R}^{N \times 1}$ denotes dimension $p = 1, \dots, P$ of the input data $\mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_P] \in \mathbb{R}^{N \times P}$. The first hidden layer is given by

$$\mathbf{h}_q^{(1)} = h\left(\sum_{p=1}^P a_{pq}^{(1)} \mathbf{x}_p + b_q^{(1)}\right), \quad q = 1, \dots, Q^{(1)}, \quad (2.1)$$

where $\mathbf{h}_q^{(1)} \in \mathbb{R}^{N \times 1}$ denotes dimension $q = 1, \dots, Q^{(1)}$ of the matrix $\mathbf{H}^{(1)} = [\mathbf{h}_1^{(1)}, \dots, \mathbf{h}_{Q^{(1)}}^{(1)}] \in \mathbb{R}^{N \times Q^{(1)}}$. In matrix-vector notations we obtain $\mathbf{H}^{(1)} = h(\mathbf{X}A^{(1)} + b^{(1)})$ with $A^{(1)} = [a_{pq}^{(1)}]_{p=1, \dots, P, q=1, \dots, Q^{(1)}}$ being a $P \times Q^{(1)}$ matrix of multipliers and $b^{(1)} = [b_1^{(1)}, \dots, b_{Q^{(1)}}^{(1)}] \in \mathbb{R}^{1 \times Q^{(1)}}$ denoting the bias vector. Here, $h(x)$ is the activation function given explicitly by $h(x) = \tanh(x)$. Similarly, the second hidden layer $\mathbf{H}^{(2)} \in \mathbb{R}^{N \times Q^{(2)}}$ is given by $\mathbf{H}^{(2)} = h(\mathbf{H}^{(1)}A^{(2)} + b^{(2)})$ where $A^{(2)}$ is a $Q^{(1)} \times Q^{(2)}$ matrix of multipliers and $b^{(2)} \in \mathbb{R}^{1 \times Q^{(2)}}$ denotes the bias vector. The output of the neural network $\mathbf{F} \in \mathbb{R}^{N \times R}$ is given by $\mathbf{F} = \mathbf{H}^{(2)}A + b$ with $A \in \mathbb{R}^{Q^{(2)} \times R}$ and $b \in \mathbb{R}^{1 \times R}$. Moreover, we assume a Gaussian noise model

$$\mathbf{y}_r = \mathbf{f}_r + \boldsymbol{\epsilon}_r, \quad r = 1, \dots, R, \quad (2.2)$$


with mutually independent $\boldsymbol{\epsilon}_r \sim \mathcal{N}(\mathbf{0}, \sigma_r^2 \mathbf{I})$. Letting $\mathbf{y} := [\mathbf{y}_1; \dots; \mathbf{y}_R] \in \mathbb{R}^{NR \times 1}$, we obtain the following likelihood

$$\mathbf{y} \sim \mathcal{N}(\mathbf{y}|\mathbf{f}, \Sigma \otimes \mathbf{I}), \quad (2.3)$$

where $\Sigma = \text{diag}(\sigma_1^2, \dots, \sigma_S^2)$ and $\mathbf{f} := [\mathbf{f}_1; \dots; \mathbf{f}_R] \in \mathbb{R}^{NR \times 1}$. One can train the parameters of the neural network by minimizing the resulting negative log likelihood.

2.1. Illustrative Example. Figure 2.1 depicts a neural network fit to a synthetic dataset generated by random perturbations of a simple one dimensional function.

3. Recurrent Neural Networks [2]. Let us consider a time series dataset of the form $\{\mathbf{y}_t : t = 1, \dots, T\}$. We can employ the following recurrent neural network



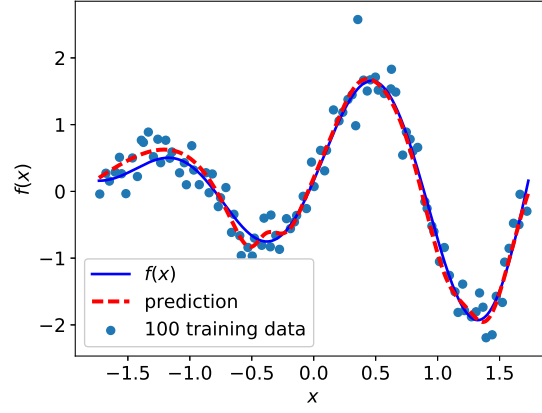


FIG. 2.1. Neural network fitting a simple one dimensional function.




FIG. 3.1. Recurrent neural network predicting the dynamics of a simple sine wave.

to model the next value $\hat{\mathbf{y}}_t$ of the variable of interest as a function of its own lagged values \mathbf{y}_{t-1} and \mathbf{y}_{t-2} ; i.e., $\hat{\mathbf{y}}_t = f(\mathbf{y}_{t-1}, \mathbf{y}_{t-2})$. Here, $\hat{\mathbf{y}}_t = \mathbf{h}_t V + c$, $\mathbf{h}_t = \tanh(\mathbf{h}_{t-1} W + \mathbf{y}_{t-1} U + b)$, $\mathbf{h}_{t-1} = \tanh(\mathbf{h}_{t-2} W + \mathbf{y}_{t-2} U + b)$, and $\mathbf{h}_{t-2} = \mathbf{0}$. The parameters U, V, W, b , and c of the recurrent neural network can be trained by minimizing the mean squared error

$$\mathcal{MSE} := \frac{1}{T-2} \sum_{t=3}^T |\mathbf{y}_t - \hat{\mathbf{y}}_t|^2.$$

3.1. Illustrative Example. Figure 3.1 depicts a recurrent neural network (with 5 lags) learning and predicting the dynamics of a simple sine wave.

4. Long Short Term Memory (LSTM) Networks [2]. A long short term memory (LSTM) network replaces the units $\mathbf{h}_t = \tanh(\mathbf{h}_{t-1} W + \mathbf{y}_{t-1} U + b)$ of a recurrent neural network with

$$\mathbf{h}_t = \mathbf{o}_t \odot \tanh(\mathbf{s}_t),$$

where

$$\mathbf{o}_t = \sigma(\mathbf{h}_{t-1} W_o + \mathbf{y}_{t-1} U_o + b_o),$$

is the output gate and

$$\mathbf{s}_t = \mathbf{f}_t \odot \mathbf{s}_{t-1} + \mathbf{i}_t \odot \tilde{\mathbf{s}}_t,$$

is the cell state. Here,

$$\tilde{\mathbf{s}}_t = \tanh(\mathbf{h}_{t-1} W_s + \mathbf{y}_{t-1} U_s + b_s).$$

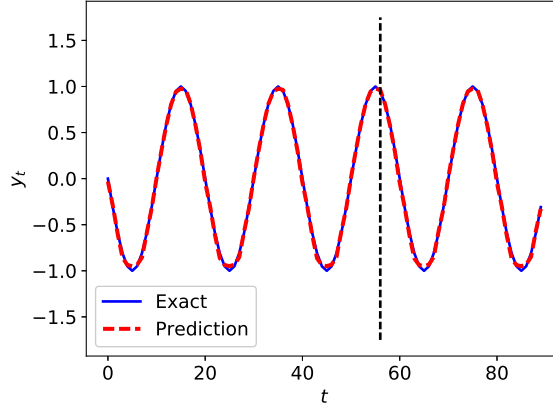


FIG. 4.1. Long short term memory network predicting the dynamics of a simple sine wave.




FIG. 4.2. Training data and samples generated by a variational auto-encoder.

Moreover,

$$\mathbf{i}_t = \sigma(\mathbf{h}_{t-1}W_i + \mathbf{y}_{t-1}U_i + b_i)$$

is the external input gate while

$$\mathbf{f}_t = \sigma(\mathbf{h}_{t-1}W_f + \mathbf{y}_{t-1}U_f + b_f),$$

is the forget gate.

4.1. Illustrative Example. Figure 4.1 depicts a long short term memory network (with 10 lags) learning and predicting the dynamics of a simple sine wave.

5. Variational Auto-encoders [1]. Let us start by the prior assumption that

$$p(\mathbf{z}) = \mathcal{N}(\mathbf{z}|\mathbf{0}, I),$$

where \mathbf{z} is a latent variable. Moreover, let us assume

$$p(\mathbf{y}|\mathbf{z}) = \mathcal{N}(\mathbf{y}|\mu_2(\mathbf{z}), \Sigma_2(\mathbf{z})),$$

where $\mu_2(\mathbf{z})$ and $\Sigma_2(\mathbf{z})$ are modeled as deep neural networks. Here, $\Sigma_2(\mathbf{z})$ is constrained to be a diagonal matrix. We are interested in minimizing the negative log likelihood $-\log p(\mathbf{y})$, where $p(\mathbf{y}) = \int p(\mathbf{y}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$. However, $-\log p(\mathbf{y})$ is not analytically tractable. To deal with this issue, one could employ a variational distribution $q(\mathbf{z}|\mathbf{y})$ and compute the following Kullback-Leibler divergence; i.e.,

$$\mathbb{KL}[q(\mathbf{z}|\mathbf{y}) || p(\mathbf{z}|\mathbf{y})] = \int \log \frac{q(\mathbf{z}|\mathbf{y})}{p(\mathbf{z}|\mathbf{y})} q(\mathbf{z}|\mathbf{y}) d\mathbf{z} = \int [\log q(\mathbf{z}|\mathbf{y}) - \log p(\mathbf{z}|\mathbf{y})] q(\mathbf{z}|\mathbf{y}) d\mathbf{z}.$$

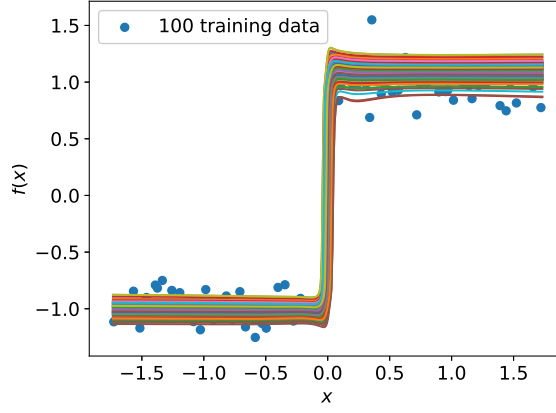


FIG. 6.1. Training data and samples generated by a conditional variational auto-encoder.

Using the Bayes rule $p(\mathbf{z}|\mathbf{y}) = \frac{p(\mathbf{y}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{y})}$, one obtains

$$\mathbb{KL}[q(\mathbf{z}|\mathbf{y}) \parallel p(\mathbf{z}|\mathbf{y})] = \int [\log q(\mathbf{z}|\mathbf{y}) - \log p(\mathbf{y}|\mathbf{z}) - \log p(\mathbf{z}) + \log p(\mathbf{y})] q(\mathbf{z}|\mathbf{y}) d\mathbf{z}.$$

Therefore,

$$\mathbb{KL}[q(\mathbf{z}|\mathbf{y}) \parallel p(\mathbf{z}|\mathbf{y})] = \log p(\mathbf{y}) + \mathbb{KL}[q(\mathbf{z}|\mathbf{y}) \parallel p(\mathbf{z})] - \int \log p(\mathbf{y}|\mathbf{z}) q(\mathbf{z}|\mathbf{y}) d\mathbf{z}.$$

Rearranging the terms yields

$$-\log p(\mathbf{y}) + \mathbb{KL}[q(\mathbf{z}|\mathbf{y}) \parallel p(\mathbf{z}|\mathbf{y})] = - \int \log p(\mathbf{y}|\mathbf{z}) q(\mathbf{z}|\mathbf{y}) d\mathbf{z} + \mathbb{KL}[q(\mathbf{z}|\mathbf{y}) \parallel p(\mathbf{z})].$$

A variational auto-encoder proceeds by minimizing the terms on the right hand side of the above equation. Moreover, let us assume that

$$q(\mathbf{z}|\mathbf{y}) = \mathcal{N}(\mathbf{z}|\mu_1(\mathbf{y}), \Sigma_1(\mathbf{y})),$$

where $\mu_1(\mathbf{y})$ and $\Sigma_1(\mathbf{y})$ are modeled as deep neural networks. Here, $\Sigma_1(\mathbf{y})$ is constrained to be a diagonal matrix. One can use

$$\mu_1(\mathbf{y}) + \epsilon \Sigma_1(\mathbf{y})^{1/2}, \quad \epsilon \sim \mathcal{N}(\mathbf{0}, I).$$

to generate samples from $q(\mathbf{z}|\mathbf{y})$.

5.1. Illustrative Example. Figure 4.2 depicts the training data and the samples generated by a variational auto-encoder.

6. Conditional Variational Auto-encoders. Conditional variational auto-encoders, rather than making the assumption that $p(\mathbf{z}) = \mathcal{N}(\mathbf{z}|\mathbf{0}, I)$, start by assuming that

$$p(\mathbf{z}|\mathbf{x}) = \mathcal{N}(\mathbf{z}|\mu_0(\mathbf{x}), \Sigma_0(\mathbf{x})),$$

where $\mu_0(\mathbf{x})$ and $\Sigma_0(\mathbf{x})$ are modeled as deep neural networks. Here, $\Sigma_0(\mathbf{x})$ is constrained to be a diagonal matrix.

6.1. Illustrative Example. Figure 6.1 depicts the training data and the samples generated by a conditional variational auto-encoder.

REFERENCES

- [1] Carl Doersch. Tutorial on variational autoencoders. *arXiv preprint arXiv:1606.05908*, 2016.
- [2] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016. <http://www.deeplearningbook.org>.